

```

import torch as t
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report

transform=transforms.Compose([transforms.ToTensor(),transforms.Normalize((0.5,),(0.5,))])
train_dataset=torchvision.datasets.MNIST(root='./data',train=True,download=True,transform=transform)
test_dataset=torchvision.datasets.MNIST(root='./data',train=False,download=True,transform=transform)

image,label=train_dataset[0]
print("Image shape:",image.shape)
print("Number of training samples:",len(train_dataset))

image,label=test_dataset[0]
print("Image shape:",image.shape)
print("Number of testing samples:",len(test_dataset))
train_loader=DataLoader(train_dataset,batch_size=32,shuffle=True)
test_loader=DataLoader(test_dataset,batch_size=32,shuffle=False)

class CNNClassifier(nn.Module):
    def __init__(self):
        super(CNNClassifier,self).__init__()
        self.conv1=nn.Conv2d(in_channels=1,out_channels=32,kernel_size=3,padding=1)
        self.conv2=nn.Conv2d(in_channels=32,out_channels=64,kernel_size=3,padding=1)
        self.conv3=nn.Conv2d(in_channels=64,out_channels=128,kernel_size=3,padding=1)
        self.pool=nn.MaxPool2d(kernel_size=2,stride=2)
        self.fc1=nn.Linear(128*3*3,128)
        self.fc2=nn.Linear(128,64)
        self.fc3=nn.Linear(64,10)

    def forward(self,x):
        x=self.pool(t.relu(self.conv1(x)))
        x=self.pool(t.relu(self.conv2(x)))
        x=self.pool(t.relu(self.conv3(x)))
        x=x.view(x.size(0),-1)
        x=nn.functional.relu(self.fc1(x))
        x=nn.functional.relu(self.fc2(x))
        x=self.fc3(x)
        return x

from torchsummary import summary
model=CNNClassifier()
if t.cuda.is_available():
    device=t.device("cuda")
    model.to(device)
print("Name: Yashwini M")
print("Reg.no: 212223230249")
summary(model,input_size=(1,28,28))
criterion=nn.CrossEntropyLoss()
optimizer=optim.Adam(model.parameters(),lr=0.001)
def train_model(model,train_loader,num_epochs):
    for epoch in range(num_epochs):
        model.train()
        running_loss=0.0
        for images,labels in train_loader:
            if t.cuda.is_available():
                images,labels=images.to(device),labels.to(device)
            optimizer.zero_grad()
            outputs=model(images)
            loss=criterion(outputs,labels)
            loss.backward()
            optimizer.step()
            running_loss+=loss.item()
        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")
print("Name: Yashwini M")
print("Reg.no: 212223230249")

train_model(model,train_loader,num_epochs=10)

def test_model(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    all_preds = []
    all_labels = []

```

```

with t.no_grad():
    for images, labels in test_loader:
        if t.cuda.is_available():
            images, labels = images.to(device), labels.to(device)

        outputs = model(images)
        _, predicted = t.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

accuracy = correct/total
print("Name: Yashwini M")
print("Register No: 212223230249")
print(f"Test Accuracy: {accuracy:.4f}")

cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(8, 6))
print("Name: Yashwini M")
print("Register No: 212223230249")
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_dataset.classes, yticklabels=test_dataset.classes)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

print("Name: Yashwini M")
print("Register No: 212223230249")
print("Classification Report:")
print(classification_report(all_labels, all_preds, target_names=[str(i) for i in range(10)]))
test_model(model, test_loader)

def predict_image(model,image_index,dataset):
    model.eval()
    image,label=dataset[image_index]
    if t.cuda.is_available():
        image=image.to(device)

    with t.no_grad():
        output=model(image.unsqueeze(0))
        _,predicted=t.max(output,1)
        class_names=[str(i) for i in range(10)]
        print("Name: Yashwini M")
        print("Reg no: 212223230249")
        plt.imshow(image.cpu().squeeze(0),cmap='gray')
        plt.title(f"Actual: {class_names[label]}\nPredicted: {class_names[predicted.item()]}")
        plt.axis("off")
        plt.show()
        print(f"Actual: {class_names[label]}\nPredicted: {class_names[predicted.item()]}")
    predict_image(model,image_index=80,dataset=test_dataset)

```

Image shape: torch.Size([1, 28, 28])
 Number of training samples: 60000
 Image shape: torch.Size([1, 28, 28])
 Number of testing samples: 10000
 Name: Yashwini M
 Reg.no: 212223230249

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 28, 28]	320
MaxPool2d-2	[-1, 32, 14, 14]	0
Conv2d-3	[-1, 64, 14, 14]	18,496
MaxPool2d-4	[-1, 64, 7, 7]	0
Conv2d-5	[-1, 128, 7, 7]	73,856
MaxPool2d-6	[-1, 128, 3, 3]	0
Linear-7	[-1, 128]	147,584
Linear-8	[-1, 64]	8,256
Linear-9	[-1, 10]	650

Total params: 249,162

Trainable params: 249,162

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.42

Params size (MB): 0.95

Estimated Total Size (MB): 1.37

Name: Yashwini M

Reg.no: 212223230249

Epoch [1/10], Loss: 0.1547

Epoch [2/10], Loss: 0.0464

Epoch [3/10], Loss: 0.0338

Epoch [4/10], Loss: 0.0262

Epoch [5/10], Loss: 0.0217

Epoch [6/10], Loss: 0.0190

Epoch [7/10], Loss: 0.0137

Epoch [8/10], Loss: 0.0141

Epoch [9/10], Loss: 0.0120

Epoch [10/10], Loss: 0.0122


Name: Yashwini M

Register No: 212223230249

Test Accuracy: 0.9917

Name: Yashwini M

Register No: 212223230249



Name: Yashwini M

Register No: 212223230249

Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	1.00	980
1	1.00	1.00	1.00	1135
2	1.00	1.00	1.00	1032
3	0.99	1.00	0.99	1010
4	1.00	0.98	0.99	982
5	0.99	0.99	0.99	892